



Week 4 Lecture 1

▼ Class	BSCCS2003
🕒 Created	@September 27, 2021 1:48 PM
🔗 Materials	
# Module #	23
▼ Type	Lecture
☰ Week #	4

Persistent Storage

Example: Gradebook

- Students: ID, name, address, ...
- Courses: ID, name, department, year, ...
- Student-Course Relationship: which students are registered for which courses

Gradebook

Aa Name	☰ IDNumber
<u>Sunil Shashi</u>	MAD001
<u>Chetana Anantha</u>	MAD002
<u>Madhur Prakash</u>	MAD003
<u>Nihal Surya</u>	MAD004
<u>Shweta Lalita</u>	MAD005
<u>Raghu Balwinder</u>	MAD006
<u>Gulshan Kuldeep</u>	MAD007
<u>Kishan Shrivatsa</u>	MAD008
<u>Purnima Sunil</u>	MAD009
<u>Nikitha Madhavi</u>	MAD010
<u>Lilavati Prabhakar</u>	MAD011
<u>Rama Yamuna</u>	MAD012

Gradebook

Aa CourseID	☰ Name
<u>EE1001</u>	Introduction to Electrical Engineering
<u>AM1100</u>	Engineering Mechanics
<u>MA1020</u>	Functions of Several Variables
<u>ME1100</u>	Thermodynamics
<u>BR1010</u>	Life Sciences

Spreadsheets

- Arbitrary data organized into rows and columns
- Operations defined on cells or ranges
- Multiple inter-linked sheets within single spreadsheet

Any kind of tabular data - expressed in tables

Relationships

- Students - Course?
- Separate entry with full details - student name, ID, address, course ID, name, department, etc.
 - Redundant
- Separate table "joining" students with courses
 - Only ID specified
 - Relation specified with "Keys"

<u>Aa</u> A	<u>≡</u> B	<u>#</u> C
<u>MAD001</u>	BT1010	78
<u>MAD002</u>	EE1001	30
<u>MAD005</u>	EE1001	68
<u>MAD009</u>	AM1100	62
<u>MAD012</u>	AM1100	77
<u>MAD007</u>	BT1010	41

Questions

- How should the underlying data should be stored?
 - Can it be made persistent - survive server restart?
- How should the relations be represented?
- Structured ways to represent, manipulate data?



Week 4 Lecture 2

▼ Class	BSCCS2003
🕒 Created	@September 27, 2021 2:41 PM
🔗 Materials	
# Module #	24
▼ Type	Lecture
☰ Week #	4

Mechanisms for Persistent Storage

In-memory data structures

```
names = ['Alice', 'Bob', 'Charlie']
courses = ['Introduction to EE', 'Applied Mech', 'Calculus']
rels = [('Alice', 'Introduction to EE'),
        ('Bob', 'Calculus'),
        ('Alice', 'Calculus'),
        ('Charlie', 'Applied Mech')]
```

- Error prone - easy to make mistakes in entry or referencing
- Does not scale
- Duplicate names?

In-memory data structures: Keys

```
names = {0: 'Alice', 1: 'Bob', 2: 'Charlie'}
courses = {0: 'Introduction to EE', 1: 'Applied Mech', 2: 'Calculus'}
rels = [(0, 0), (1, 2), (0, 2), (2, 1)]
```

- Data entry errors less likely
- Duplicates not a problem - Unique key

Objects

```
class Student:
    id_next = 0 # Class variable
    def __init__(self, name):
        self.name = name
        self.id = Student.id_next
        Student.id_next += 1
```

- Auto-initialize ID to ensure unique
- Functions to set/get values

```
class Student:
    id_next = 0 # Class variable
    def __init__(self, name, hostel):
        self.name = name
        self.id = Student.id_next
        self.hostel = hostel
        Student.id_next += 1
```

- Add a new field to the object easily

But what about persistence?

- In-memory data structures are lost when the program is killed, server is shut down or rebooted (restarted)
- Save to disk? Structured data?
 - Python Pickle and similar modules
 - CSV - Comma Separated Values
 - TSV - Tab Separated Values
- Essentially the same as spreadsheets - limited flexibility

Spreadsheets

- Naturally represent tabular data
- Extension, adding fields are easy
- Separate sheets for relationships

Problems with spreadsheets:

- Lookups, cross-referencing harder than dedicated database
- Stored procedures - limited functionality
- Atomic operations - no clear definition

Relational Databases - SQL

Shameless plug, I do have notes for DBMS here

- From IBM ~ 1970s
- Data is stored in tabular format:
 - Columns of tables: fields (name, address, department, ...)
 - Rows of tables: individual entries (student1, student2, ...)

Unstructured Databases - NoSQL

- Easily add/change fields
- Arbitrary data
- NoSQL
 - MongoDB
 - CouchDB
 - ...

- Flexible, but potential loss of validation



Week 4 Lecture 3

▼ Class	BSCCS2003
🕒 Created	@September 27, 2021 3:39 PM
🔗 Materials	
# Module #	25
▼ Type	Lecture
☰ Week #	4


Relations and ER diagram

Relationship types

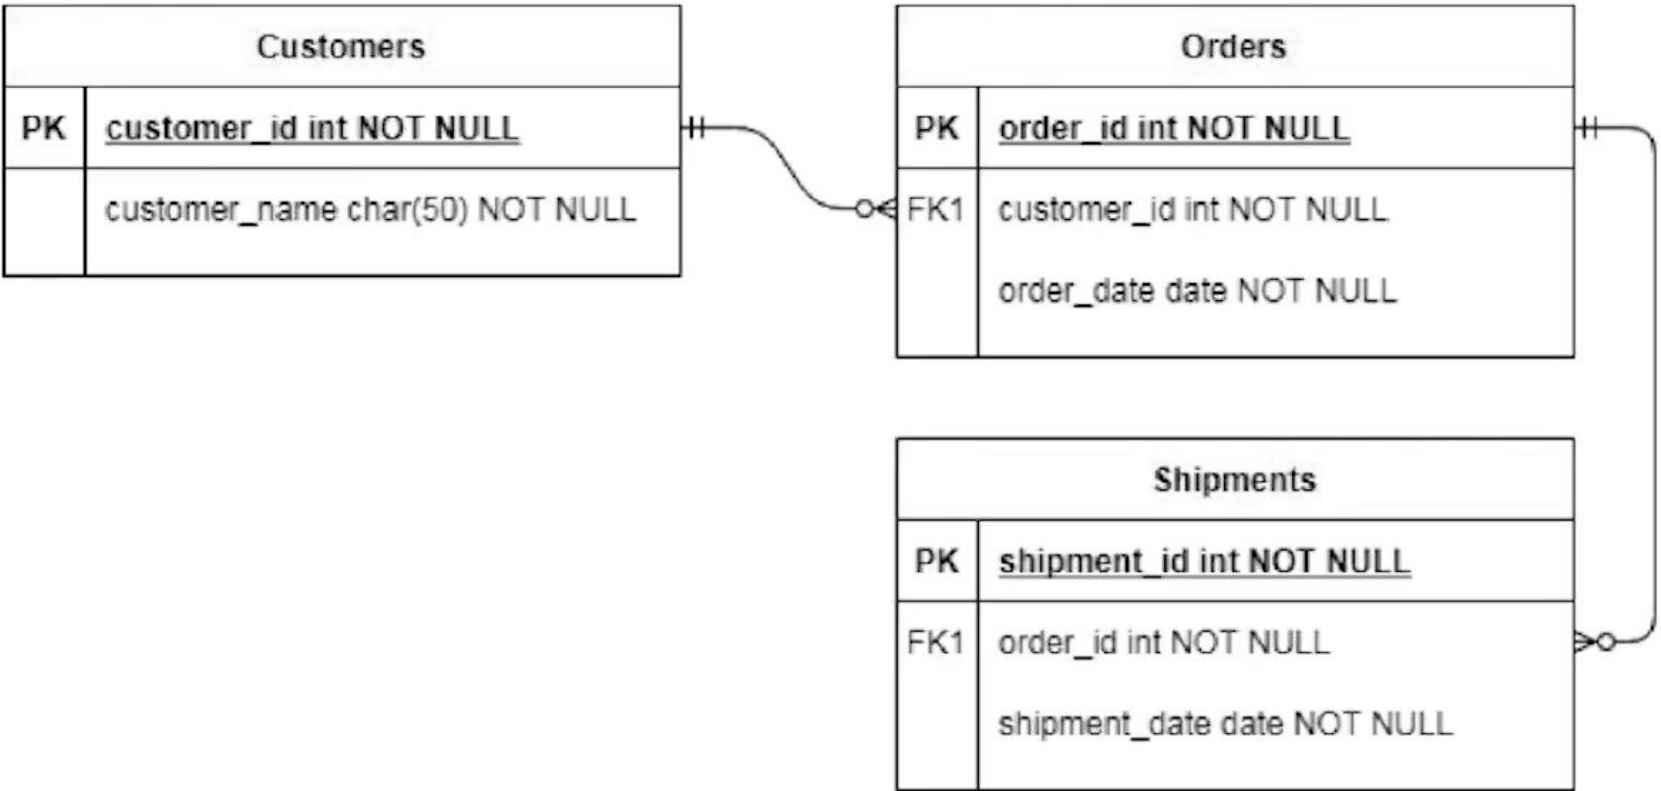
- One to One:
 - One student has one roll number
 - One roll number uniquely identifies one student
 - **Example:** assign unique message-ID to each email in the inbox
- One to Many (Many to One):
 - One stays in only one hostel
 - One hostel has many students
 - **Example:** save emails in folders - one email is in only one folder
- Many to Many:
 - One student can register for many courses
 - One course be taken by many students
 - **Example:** assign labels to emails - one email can have many labels and vice-versa

Diagrams




- Entity-Relationship (ER)
- Unified Modeling Language (UML)
- Class relation



Entity-Relationship Diagram



ER examples



Tool: Draw.io → <https://app.diagrams.net>



Week 4 Lecture 4

▼ Class	BSCCS2003
🕒 Created	@September 27, 2021 4:29 PM
🔗 Materials	
# Module #	26
▼ Type	Lecture
☰ Week #	4

SQL

Relational Databases - SQL

- From IBM ~ 1970s
- Data is stored in tabular format:
 - Columns of tables: fields (name, address, department, ...)
 - Rows of tables: individual entries (student1, student2, ...)
- Key: unique way of accessing a given row
 - **Primary Key:** important for fast access to large databases
 - **Foreign Key:** connect to a different table - Relationships

Queries

- Retrieve data from a database

eg: "Find students with name beginning with A"
"Find all courses offered in 2021"

Structured Query Language

- English-like, but structured
- Quite verbose
- Specific mathematical operations:
 - Inner join

- Outer join

Example: Inner join

<u>Aa</u> Name	<u>≡</u> IDNumber	<u>#</u> hostelID
<u>Sunil Shashi</u>	MAD001	1
<u>Chetana Anantha</u>	MAD002	2
<u>Madhur Prakash</u>	MAD003	2
<u>Nihal Surya</u>	MAD004	3
<u>Shweta Lalita</u>	MAD005	2
<u>Raghu Balwinder</u>	MAD006	3
<u>Gulshan Kuldeep</u>	MAD007	1
<u>Kishan Shrivatsa</u>	MAD008	1
<u>Purnima Sunil</u>	MAD009	2
<u>Nikitha Madhavi</u>	MAD010	1
<u>Lilavati Prabhakar</u>	MAD011	3
<u>Rama Yamuna</u>	MAD012	3

<u>Aa</u> ID	<u>≡</u> Name	<u>#</u> Capacity
<u>1</u>	Jamuna	300
<u>2</u>	Ganga	300
<u>3</u>	Brahmaputra	500

Student - Hostel mapping

```
SELECT Students.name, Hostels.name
FROM Students
INNER JOIN Hostels
ON Students.hostelID = Hostels.ID;
```

---- OUTPUT ----

Sunil Shashi, Jamuna
Chetana Anantha, Ganga

Cartesian product

- N entries in table 1
- M entries in table 2
- $M \times N$ combinations - filter on them

Powerful SQL queries can be constructed

Example: find all the students in Calculus

- Find ID number for the course
- Look up StudentsCourses table to find all entries with this course ID
- Look up Students to find the names of students with these IDs

```
SELECT s.name
FROM Student s
JOIN StudentsCourses sc ON s.IDNumber = sc.studentID
JOIN Courses c ON c.ID = sc.courseID
WHERE c.name = 'Calculus';
```

Summary

- Models - persistent data storage
- Mechanisms:
 - CSV
 - spreadsheets
 - SQL
 - NoSQL

- Entities and Relationships
 - Different ways of representing

No details on display, views or what kind of updates permitted