



Week 7 Lecture 1

▼ Class	BSCCS2003
🕒 Created	@October 18, 2021 10:50 AM
🔗 Materials	
# Module #	38
▼ Type	Lecture
# Week #	7

Memory Hierarchy

Types of Storage Elements

- On-chip registers: 10s-100s of bytes
- SRAM (Cache): 0.1-1MB # Static RAM
- DRAM (Memory/RAM of a PC): 0.1-16GB
- Solid-State Disk (SSD) - Flash: 1-128GB
- Magnetic Disk (HDD-Hard Disk Drive): 0.1-10TB

Storage Parameters

- **Latency:** Time to read the first value from a storage location (lower is better)
 - Register < SRAM < DRAM < SSD < HDD
 - So, the turnaround time when trying to read something from a ...
 - Register: in order of nanoseconds
 - SRAM: tens to hundreds of nanoseconds
 - DRAM: several microseconds
 - SSDs: hundreds of microseconds
 - HDDS: order of milliseconds
- **Throughput:** number of bytes/second that can be read (higher is better)
 - DRAM > SSD > HDD (registers, SRAM have limited capacity)

- **Density:** Number of bits stored per unit area/cost (higher is better)
 - Volume manufacture important
 - HDD > SSD > DRAM > SRAM > Registers

Computer Organization

- CPU has as many registers as possible
- Backed by L1, L2, L3 cache (SRAM)
- Backed by several Gigabytes (GBs) of DRAM working memory
- Backed by SSD for high throughput
- Backed by HDD for high capacity
- Backed by long-term storage, backup

Cold Storage

- Backups & archives:
 - Huge amounts of data
 - Not read very often
 - Can tolerate high read latency
- Amazon Glacier, Google, Azure Cold/Archive storage classes
- High latency of retrieval → up to 48 hours
- Very high durability
- Very low cost

Impact on Application Development

- Plan the storage needs based on the application growth
- Speed of app determined by types of data stored, how stored
- Some data stores are more efficient for some types of read/write operations

Developer must be aware of the choices and what kind of DB to choose for a given application



Week 7 Lecture 2

▼ Class	BSCCS2003
🕒 Created	@October 18, 2021 2:13 PM
🔗 Materials	
# Module #	39
▼ Type	Lecture
# Week #	7

Data Search

O() notation - The Big O notation

- Used in study of algorithmic complexity
- Rough approximation: "order of magnitude", "approximately", etc.
- Main concepts:
 - $O(1)$ - constant time, independent of input size: Excellent!
 - $O(\log N)$ - logarithmic in input size - grows very slowly with input: Very good
 - $O(N)$ - linear in input size - often the baseline - Would like to do better
 - $O(N^k)$ - polynomial (quadratic, cubic, etc.) - Not good as the input size grows
 - $O(k^N)$ - exponential - **VERY BAD**: Won't even work for reasonably small inputs

Searching for an element in the memory

Unsorted data in a linked list (or an array)

- Start from the beginning
- Proceed stepwise, comparing each element
- Stop if found and return the LOCATION
- If end-of-list, return NOTFOUND

$O(N)$

Sorted data in an array

- Look at the middle element in an array:
 - Greater than the target?
 - Search in the lower half
 - Lesser than the target?
 - Search in the upper half
- Switch focus to the new array: Half the original size
 - Repeat

$O(\log N)$: *Binary Search*

Problems with arrays

- Size must be fixed ahead of time
- Adding new entries requires re-sizing - can try oversize, but eventually ...
- Maintaining sorted order $O(N)$:
 - Find the location to insert
 - Move all further elements by 1 to create a gap
 - Insert
- Deleting
 - Find the location, delete
 - Move all the entries down by 1 step

Alternatives

- Binary Search Tree
 - Maintaining sorted order is easier: growth of the tree
- Self-Balancing
 - BST can easily tilt to one side and grow downwards
 - Red-black, AVL, B-tree ... more complex, but still reasonable
- Hash tables
 - Compute an index for an element: $O(1)$
 - Hope the index for each element is unique!
 - Difficult but is doable in many cases



Week 7 Lecture 3

▼ Class	BSCCS2003
🕒 Created	@October 18, 2021 2:42 PM
🔗 Materials	
# Module #	40
▼ Type	Lecture
# Week #	7

Database Search

Databases (Tabular)

- Tabular with many columns
- Want to search quickly on some columns
- Maintain "INDEX" of columns to search on
 - Store a sorted version of column
 - Needs column to be "comparable": integer, short string, date/time, etc.
 - Long text fields are not good for index
 - Binary data not good as well

Example: MySQL

8.3.9 Comparison of B-Tree and Hash Indexes

Understanding the B-tree and hash data structures can help predict how different queries perform on different storage engines that use these data structures in their indexes, particularly for the MEMORY storage engine that lets you choose B-tree or hash indexes.

- [B-Tree Index Characteristics](#)
- [Hash Index Characteristics](#)

Source: <https://dev.mysql.com/doc/refman/8.0/en/index-btree-hash.html>

Index-friendly Query

```
SELECT * FROM tbl_name WHERE key_col LIKE 'Patrick%';
SELECT * FROM tbl_name WHERE key_col LIKE 'Pat%ck%';
```

In the first statement, only rows with `'Patrick' <= key_col < 'Patricl'` are considered

In the second statement, only rows with `'Pat' <= key_col < 'Pau'` are considered

Index-unfriendly Query

```
SELECT * FROM tbl_name WHERE key_col LIKE '%Patrick%';
SELECT * FROM tbl_name WHERE key_col LIKE other_col;
```

In the first statement, the `LIKE` value begins with a wildcard character

In the second statement, the `LIKE` value is not a constant

Multi-column index

- (index_1, index_2, index_3): compound index on 3 columns:
 - first sorted on index_1, then on index_2, then on index_3
 - all values with same index_1 will be sorted on index_2
 - all values with same index_1 and index_2 will be sorted on index_3
- Example: (date-of-birth, city-of-birth, name)
 - can query for all the people born on the same date in the same city with the same name easily

Multi-index friendly

```
... WHERE index_part1=1 AND index_part2=2 AND other_column=3

/* index = 1 OR index = 2 */
... WHERE index=1 OR A=10 AND index=2

/* optimized like "index_part1='hello'" */
... WHERE index_part1='hello' AND index_part3=5

/* use index on index1 but not on index2 or index3 */
... WHERE index1=1 AND index2=2 OR index1=3 AND index3=3;
```

Multi-index unfriendly

```
/* index part1 is not used */
... WHERE index_part2=1 AND index_part3=2

/* Index not used in both parts of the WHERE clause */
... WHERE index=1 OR A=10

/* No index spans all rows */
... WHERE index_part1=1 OR index_part2=10
```

Hash-index

- Only used in in-memory tables
- Only for equality comparisons - cannot handle ranges
- Does not help with "ORDER BY"
- Partial key prefix cannot be used
- But VERY fast where applicable

Query Optimization

Database specific:

- <https://dev.mysql.com/doc/refman/8.0/en/index-btree-hash.html>
- <https://www.sqlite.org/optoverview.html>
- <https://www.postgresql.org/docs/12/geqo.html>

Summary

- Setting up queries properly impacts application performance
- Building proper indexes is crucial to good search
- Compound indexes, multiple indexes etc. possible
 - Too many can be a waste of space
- Make use of structure in data to organize it properly



Week 7 Lecture 4

▼ Class	BSCCS2003
🕒 Created	@October 18, 2021 3:42 PM
🔗 Materials	
# Module #	41
▼ Type	Lecture
# Week #	7

SQL v/s. NoSQL

SQL

- Structured **Q**uery Language
 - Used to query DB that have structure
 - Could also be used for CSV files, spreadsheets, etc.
- Closely tied to RDBMS - Relational DBs
 - COLUMNS/Fields
 - Tables of data hold relationships
 - All entries in a table **must** have same set of columns
- Tabular DBs
 - Efficient indexing possible - use specified columns
 - Storage efficiency: prior knowledge of data size

Problems with tabular DBs

- Structure (good? bad?)
- All rows in a table must have same set of columns

Example:

- Student - hostel ⇒ mess
- Student - day-scholar ⇒ gate pass for vehicle

- Table? Column for mess, column for gate pass?

Alternate ways to store: Document DBs

- Free-form (unstructured) documents
 - Typically JSON encoded
 - Still structured, but each document has own structure
- **Examples:**
 - MongoDB
 - Amazon DocumentDB

```
[
  {
    "year": 2013,
    "title": "Turn It Down, Or Else!",
    "info": {
      "directors": ["Alice Smith", "Bob Jones"],
      "release_date": "2013-01-18T00:00:00Z",
      "rating": 6.2,
      "genres": ["Comedy", "Drama"],
      "image_url": "http://ia.media-imdb.com/images/N/09ERWAU7FS797AJ7LU8HN",
      "plot": "A rock band plays their music at high volumes, annoying the neighbours",
      "actors": ["David Matthewman", "Jonathan G. Neff"]
    }
  },
  {
    "year": 2015,
    "title": "The Big New Movie",
    "info": {
      "plot": "Nothing happens at all",
      "rating": 0
    }
  }
]
```

Alternate ways to store: Key-Value

- Python dictionary, C++ OrderedMap, etc. → dictionary/hash table
- Map a key to a value
- Store using search trees or hash tables
- Very efficient key lookup, not good for range type queries
- **Example:**
 - Redis
 - BerkeleyDB
 - memcached
- Often used alongside other DBs for "in-memory" fast queries

Alternate ways to store: Column store

- Traditional relational DBs store all values of a row together on the disk
 - Retrieving all entries of a given row is very fast
- Instead store all entries in a column together
 - Retrieve all values of a given attribute (age, place of birth, ...) very fast
- **Example:**
 - Cassandra
 - HBase

Alternate ways to store: Graphs

- Friend-of-a-friend, social networks, maps: graph oriented relationships
- Different degrees (number of out-going edges), weights of edges, nodes, etc.

- Path-finding more important than just search
 - Connections, knowledge discovery
- Examples:
 - Neo4j
 - Amazon Neptune

Alternate ways to store: Time Series DBs

- Very application specific: Store some metric or values as a function of time
- Used for log analysis, performance analysis, monitoring
- Queries:
 - How many hits between T1 and T2?
 - Avg. number of requests per second?
 - Country from where maximum requests came in the past 7 days?
- Typical RDBMS completely unsuitable - same for most alternatives
- **Example:**
 - RRDTool
 - InfluxDB
 - Prometheus
 - Elasticsearch
 - grafana

NoSQL

- Started out as an "alternative" to SQL
- But SQL is just a query language - can be adapted for any kind of query, including from a document store or graph
- "Not-only-SQL"
- Additional query patterns for other types of data stores

ACID

- **Transaction:** core principle of DB
- **ACID:**
 - Atomic
 - Consistent
 - Isolated
 - Durable
- Many NoSQL DBs sacrifice some part of ACID
 - **Example:** Eventual consistency instead of consistency, for performance
- But, there can be ACID compliant NoSQL DB as well

Why not ACID?

- Consistency hard to meet: especially when scaling/distributing
- Eventual consistency easier to meet
- Example:
 - A (located in India) and B (located in USA) both add C as a friend on Facebook
 - Order of adding does not matter
 - Temporarily seeing C in A's list but not B, or B's list but not A - not a catastrophe

- Financial transactions **absolutely require ACID**
 - Consistency is paramount: even a split second of inconsistent data can cause problems

A word on storage

- In-memory:
 - Fast
 - Does not scale across machines
- Disk
 - Different data structures, organization needed



Week 7 Lecture 5

▼ Class	BSCCS2003
🕒 Created	@October 18, 2021 6:50 PM
🔗 Materials	
# Module #	42
▼ Type	Lecture
# Week #	7

Scaling

Replication and Redundancy

- Redundancy:
 - Multiple copies of same data
 - Often used in connection with backups - even if one fails, others survive
 - One copy is still the master
- Replication:
 - Usually in context of performance
 - May not be for purpose of backup
 - Multiple sources of same data - less chance of server overload
 - Live replication requires careful design

BASE vs ACID

- "Basically Available", "Soft state", "Eventually consistent"
 - Winner of worst acronym award, lol
- Eventual consistency instead of Consistency
 - Replicas can take time to reach a consistent state
- Stress on high availability of data

Replication in traditional DBs

- RDBMS replication is possible
- Usually server cluster in the same data center
 - Load balancing
- Geographically distributed replication harder
 - Latency constraints for consistency

Scale-up vs Scale-out

- **Scale-up:** Traditional approach
 - Larger machine
 - More RAM
 - Faster network, processor
 - Requires machine restart with each scale change
- **Scale-out:**
 - Multiple servers
 - Harder to enforce consistency etc. → better suited for NoSQL/non-ACID
 - Better suited to cloud model: Google, AWS etc. to provide automatic scale-out, cannot do auto scale-up

Application Specific

- Financial transactions:
 - Cannot afford even slightest inconsistency
 - Only scale-up possible
- Typical web application
 - Social networks, media: Eventual consistency is OK
 - e-Commerce: only the financial part needs to go to ACID



Week 7 Lecture 6

▼ Class	BSCCS2003
🕒 Created	@October 18, 2021 7:05 PM
🔗 Materials	
# Module #	43
▼ Type	Lecture
# Week #	7

Security of Databases

SQL in context of an application

- Non-MVC app: can have direct SQL queries anywhere
- MVC: only in controller, but any controller can trigger a DB query

So, what is dangerous about queries?

Typical HTML form

```
<form>
  Username:
  <input type="text" name="name">
  <br />
  Password:
  <input type="password" name="password">
  <br />
</form>
```

Username:

Password:

Code

```
name = form.request['name']
pswd = form.request['password']

sql = 'SELECT * FROM Users WHERE Name = "' + name + '" AND Pass = "' + pswd + "'"
```

Example Input vs SQL

Username: *abcd*

Password: *pass*

```
SELECT * FROM Users WHERE Name = "abcd" AND Pass = "pass"
```

Username: " or ""="

Password: " or ""="

```
SELECT * FROM Users WHERE Name = "" or "" AND Pass = "" or ""
```

→ *The above query checks for, basically, nothing and returns everything from the table* **Users**

```
sql = "SELECT * FROM Users WHERE Name = " + name
```

Input:

```
a; DROP TABLE Users;
```

Query:

```
SELECT * FROM Users WHERE Name = a; DROP TABLE Users;
```

Problem

- Parameters from the HTML taken without validation
- Validation:
 - Are they valid text data (no special characters, other symbols)
 - No punctuation or other invalid input
 - Are they the right kind of input (text numbers, email, date)?
- Validation **MUST** be done before the DB query - even if you have validation in the HTML or JavaScript - not good enough
 - Direct HTTP requests can be made with junk data

Web Application Security

- SQL injection
 - Use known frameworks, best practices, validation
- Buffer overflows, input overflows
 - Length of inputs, queries
- Server level issues - protocol implementation?
 - Use known servers with good track record of security
 - Update all patches
- Possible outcomes:
 - loss of data - deletion
 - exposure of data - sensitive information leak
 - manipulation of data - change

HTTPS

- Secure sockets: Secure communication between client and server
- Server certificate:
 - based on DNS: Has been verified by some trusted 3rd party
 - difficult to spoof
 - based on Mathematical properties - ensure very low probability of mistakes match

- However:
 - Only secures the link for data transfer - does not perform validation or safety checks
 - Negative impact on "caching" of resources like static files
 - Some overhead on performance

Summary

- Internet and Web security are complex: Enough for a course in themselves
- Generally recommended to use known frameworks with trusted track records
- Code audits
- Patch updates on OS, server, network stack, etc. essential

App developers should be very careful of their code, but also aware of problems on other levels of stack.